

RESEARCH ARTICLE OPEN ACCESS

A Novel Predictive Temporal Bridge Architecture for Ground-to-Spacecraft Command Synchronization Under Interplanetary Light-Time Delay

Jason Pandian¹  | I. Kala² 

¹Department of Information Technology, Nehru Institute of Technology, Coimbatore, Tamil Nadu, India | ²Department of Computer Science and Engineering, PSG Institute of Technology and Applied Research, Coimbatore, Tamil Nadu, India |

Correspondence: Jason Pandian (pandiajason@gmail.com)

Received: 02 June 2026 | Revised: – | Accepted: –

Funding: This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Keywords: deep space communication | predictive telemetry | delay-tolerant networking | ASIC co-processor | reality reconciliation | orbital insertion | human-machine interface | situational awareness

ABSTRACT

Electromagnetic signal propagation delay between Earth and crewed interplanetary spacecraft constitutes a fundamental constraint on command integrity and mission safety. For a crewed Mars transit at median opposition, the one-way light-time reaches $m \approx 240$ s, imposing a minimum round-trip command latency of $2m \approx 480$ s. Under such conditions, conventional reactive ground control is operationally non-viable: by the time Earth receives a telemetry anomaly and issues a correction command, the spacecraft state has evolved $2m$ seconds beyond the point of intervention, invalidating the command's physical assumptions. This paper proposes and formally evaluates the *Predictive Temporal Bridge* (PTB)—a dual-AI co-processor architecture that eliminates effective command latency by projecting the spacecraft state $2m$ seconds forward in time at the point of ground command generation. The Earth Predictive Engine computes $\hat{\mathcal{J}}(t + 2m)$ from the received telemetry packet $\mathcal{B}(t)$ using a physics-based integration of orbital mechanics, thermal dissipation models, and thruster erosion profiles. The resulting preemptive correction command $\mathcal{C}$ arrives at Mars precisely synchronized with the local spacecraft clock, where the Cabin Reconciliation Engine validates it against live sensor truth via per-channel deviation analysis $\Delta_k = T_k(t + 2m) - \hat{T}_k(t + 2m)$. A Finite State Machine (FSM) arbitrates command authorization, blocking execution if any Δ_k exceeds the safety threshold δ_{safe} , preserving crew autonomy and preventing automation-induced failures. The PTB is physically realized as a low-power Predictive Temporal Processing Unit (PTPU) ASIC interfacing with the spacecraft avionics via SpaceWire and MIL-STD-1553 buses. Validation is performed via a high-fidelity software simulation of a Mars periapsis-raising orbital insertion burn incorporating a progressive gimbal seal thermal degradation anomaly. Quantitative results confirm a mean absolute prediction error of $|\overline{\Delta}| = 1.1$ °C on thermal channels and $|\overline{\Delta}| = 0.003\%$ on trajectory deviation, both well within the safe operating envelope, and an effective reduction in perceived command latency to near-zero—a $> 99\%$ improvement over unassisted round-trip protocols.

1 | Introduction

Every crewed deep-space mission must confront a communication constraint that is, unlike most engineering problems, non-negotiable: the finite speed of light. At Mars opposition, electromagnetic signals require between 182 and 1,342 seconds one-way, with a median value near 240 seconds [1]. This delay, denoted m ,

means that any question posed by a Mars crew member will not reach Earth for m seconds, and any Earth response will take another m seconds to return—a total round-trip latency of $2m \approx 8$ minutes at median separation, growing up to 45 minutes at maximum distance.

The consequences cascade far beyond simple system performance; they disrupt the core social fabric of human collaboration.

In human-to-human systems under extreme stress, relational maintenance and mutual responsiveness are essential [2]. When latency breaks the synchronous loop, conversational flows degenerate into isolated monologue transmissions. The psychological impact on the crew is a severe sense of temporal displacement and isolation from the human species [3]. [4] demonstrated in analog Mars missions that communication delays up to 22 minutes measurably elevate Mission Controller workload, anxiety, and stress, with error rates doubling in off-nominal scenarios. Conversely, [5] showed that integrating human-in-the-loop AI can alleviate performance degradation, but did not address the relational and social breakdowns that occur under delayed conditions.

Traditional space engineering has responded to this challenge by advocating for unilateral crew autonomy, pushing onboard systems to make the crew entirely independent of ground support. While technically logical, this "autonomy bubble" induces profound relational decay: crews begin to perceive Earth Mission Control as progressively out of touch, leading to friction, psychological reactance, and mission-threatening isolation [3]. Redesigning the *computer-mediated interface* to create an adaptive temporal architecture is an urgent, human-centered alternative. By using dual-AI predictive modeling to construct a virtual "present," we can bypass the cognitive and relational stutter of multi-minute lag.

This paper introduces the **Predictive Temporal Bridge (PTB)**, an AI-mediated communication (AI-MC) architecture designed to overcome this temporal distance. Following the foundational agenda of [6], we define AI-MC as communication where an AI agent acts on behalf of or transforms the communicative exchanges between human actors. The PTB represents a radical expansion of AI-MC: rather than modifying verbal content or generating automated text suggestions, the AI synchronizes the temporal frame of reference of the human interlocutors. The human commander on Mars receives ground commands that feel instantaneous and perfectly synchronized with their immediate physical state, restoring a feeling of synchronous presence.

2 | Related Works

2.1 | Delay-Tolerant Networking and Interplanetary Routing

Interplanetary communication architectures have long been studied through the lens of Delay-Tolerant Networking (DTN), standardized in IETF RFC 5050, which employs store-and-forward bundle routing to accommodate long and variable propagation delays. [7] extended the DTN paradigm by integrating machine learning classifiers to dynamically predict link availability and pre-route bundles across hybrid satellite relay constellations, reducing effective end-to-end delivery latency. Similarly, [8] demonstrated that cognitive radio techniques applied to deep-space transponders can autonomously optimize modulation schemes and power spectral density in response to real-time link margin estimates, maximizing throughput across the DSN. While these approaches address physical-layer and network-layer efficiency, they do not resolve the fundamental problem of *command temporal mismatch*: the fact that a command generated at Earth arrival time $(t + m)$ is physically invalid by the time it reaches the spacecraft at $(t + 2m)$, because the spacecraft state has continued to evolve.

2.2 | Predictive Teleoperation and Forward-Projection Displays

The concept of masking round-trip delay through state prediction has well-established roots in space teleoperation. The Smith Predictor [9] demonstrated that a process model running in parallel with a delayed plant can synthesize a predictive error signal that partially compensates for the destabilizing effect of transport delay in closed-loop control systems. NASA JPL's "phantom robot" [10] operationalized this concept visually: a simulated kinematic overlay of a robotic manipulator was rendered in real-time for the operator, allowing motor command confirmation without waiting for the physical round-trip video signal. Within aerospace hardware, [11] extended predictive edge AI to beam-jitter compensation in deep-space laser communication links, demonstrating sub-millisecond actuator latency correction using onboard neural inference. These works collectively validate physics-based forward projection as a viable technique for latency compensation in space systems. However, they address single-operator, single-machine teleoperation scenarios and do not provide a framework for *bidirectional state synchronization* between two operational teams (ground and crew) across the full two-way light-time barrier.

2.3 | Human-in-the-Loop AI for Mission Operations

Recent work has begun bridging the gap between network-level latency mitigation and human-factor considerations. [5] proposed a human-in-the-loop AI framework for space communication delay mitigation, demonstrating that AI-mediated task decomposition can reduce operator workload under delayed feedback conditions. [4] provided empirical evidence from analog Mars missions that communication delays of 22 minutes produce decision error rates exceeding twice the nominal baseline, establishing a quantitative performance floor that fully autonomous systems must surpass. The PTB architecture presented in this paper addresses the gap left by these works by introducing a formal, mathematically grounded *dual-AI command synchronization protocol* in which both endpoints maintain coupled predictive models, and command validity is gated by a rigorous per-channel reconciliation threshold before any actuation authority is transferred to the crew.

3 | Theoretical Framework

3.1 | Temporal Command Validity and the Latency Mismatch Problem

Let $\mathcal{T}(t)$ denote the full spacecraft state vector at mission time t , comprising orbital position $\mathbf{r}(t)$, velocity $\mathbf{v}(t)$, attitude quaternion $\mathbf{q}(t)$, and onboard sensor telemetry channels. When a telemetry packet is transmitted at time t and received by Earth at time $t + m$, the ground station observes a stale state $\mathcal{T}(t)$ while the spacecraft has evolved to $\mathcal{T}(t + m)$. Any correction command $\mathcal{C}$ generated at $t + m$ and transmitted back to Mars arrives at spacecraft time $t + 2m$, when the true state is $\mathcal{T}(t + 2m)$. The *temporal command mismatch error* is therefore:

$$\epsilon_{\text{mismatch}}(t) = \mathcal{T}(t + 2m) - \mathcal{T}(t), \quad (1)$$

which grows monotonically with m and with the rate of change of the spacecraft dynamics. For a spacecraft undergoing an active propulsion burn, $\|\epsilon_{\text{mismatch}}\|$ can reach safety-critical magnitudes within the round-trip window, invalidating commands computed against the original state $\mathcal{T}(t)$.

3.2 | Predictive State Estimation for Command Compensation

The PTB mitigates Eq. (1) by computing a forward-projected state estimate $\hat{\mathcal{T}}(t + 2m)$ at the time of command generation. The Earth Predictive Engine numerically integrates the spacecraft equations of motion from the received state:

$$\hat{\mathcal{T}}(t + 2m) = \mathcal{T}(t) + \int_t^{t+2m} f(\mathcal{T}(\tau), \vec{I}(t)) d\tau + \epsilon_p, \quad (2)$$

where $f(\cdot)$ encodes orbital mechanics (Cowell's method with J2 perturbation), thermal dissipation models, and thruster erosion profiles, $\vec{I}(t)$ is the crew intent vector received in the telemetry packet, and $\epsilon_p \sim \mathcal{N}(0, \Sigma_p)$ represents the bounded prediction uncertainty. The resulting command $\mathcal{C}$ is computed against $\hat{\mathcal{T}}(t + 2m)$ rather than the stale $\mathcal{T}(t)$, achieving near-zero effective temporal mismatch when $\|\epsilon_p\| < \delta_{\text{safe}}$.

3.3 | Reconciliation Gating and Automation Safety

A critical requirement for flight-certifiable AI-assisted command systems is the explicit prevention of automation bias—the tendency of human operators to defer to automated recommendations without adequate scrutiny [6]. The PTB addresses this through a formal *dual-gate reconciliation protocol*. Upon delivery of $\mathcal{C}$ at time $t + 2m$, the Cabin Reconciliation Engine computes per-channel deviations:

$$\Delta_k = T_k(t + 2m) - \hat{T}_k(t + 2m), \quad k \in \{1, \dots, N_{\text{ch}}\}, \quad (3)$$

and evaluates the composite safety condition:

$$\text{AUTH} = \bigwedge_{k=1}^{N_{\text{ch}}} (|\Delta_k| \leq \delta_{\text{safe},k}). \quad (4)$$

Only when $\text{AUTH} = \text{TRUE}$ is the command presented to the crew for voluntary execution. If any channel fails the threshold ($\text{AUTH} = \text{FALSE}$), the FSM transitions to a DELTA-HIGH blocking state and requires explicit crew override, ensuring operational accountability remains with the human commander at all times [12].

3.4 | Distributed Situational Awareness Synchronization

[12] formalizes situational awareness (SA) as perception (Level 1), comprehension (Level 2), and projection (Level 3) of environmental states. In distributed mission operations, both the ground team and the crew must maintain *compatible Level-3 SA projections*—i.e., their models of the spacecraft's future state must agree to within acceptable bounds [13]. Without prediction, Earth's Level-3 SA is anchored at $\mathcal{T}(t)$ while the crew projects from $\mathcal{T}(t +$

$m)$, creating an irreconcilable m -second SA divergence. The PTB closes this gap by ensuring Earth's Level-3 projection is $\hat{\mathcal{T}}(t + 2m)$ and the crew's reality at command delivery is $\mathcal{T}(t + 2m)$, with divergence bounded by the reconciliation threshold δ_{safe} .

4 | The Predictive Temporal Bridge Architecture

The PTB consists of three distinct components: the Mars Broadcast Module, the Earth Predictive Engine, and the Cabin Reconciliation Engine. The overall architecture is modeled as a dynamic time-synchronized feedback loop (illustrated in Figure 1).




FIGURE 1 | System Architecture of the Predictive Temporal Bridge. The Earth AI forward-projects the ship's state by $2m$ seconds, generating preemptive commands that arrive precisely synchronized with the spacecraft's local time $t + 2m$.

4.1 | Mars Broadcast Module

At spacecraft time t , the Cabin AI continuously packages the physical state $\mathcal{T}(t)$ along with the crew's operational intent vector $\vec{I}(t)$ and system noise envelopes $\Sigma(t)$ into a broadcast packet $\mathcal{B}(t)$:

$$\mathcal{B}(t) = \{\mathcal{T}(t), \vec{I}(t), \Sigma(t)\}. \quad (5)$$

The intent vector $\vec{I}(t)$ is a high-fidelity representation of the crew's current actions, planned thruster sequences, and behavioral metrics, providing the Earth AI with the contextual constraints necessary for accurate forward projection.

4.2 | Earth Predictive Engine

Upon receiving $\mathcal{B}(t)$ at Earth clock $t + m$, the Earth AI immediately projects the spacecraft state forward by $2m$ seconds to estimate its state at the exact moment the ground command will arrive back at Mars:

$$\hat{\mathcal{T}}(t + 2m) = \mathcal{T}(t) + \int_t^{t+2m} f(\mathcal{T}(\tau), \vec{I}(t)) d\tau + \epsilon, \quad (6)$$

where $f(\cdot)$ represents the physics-based system dynamic equations (including orbital mechanics, thermal dissipation, and thruster erosion profiles), and ϵ represents the uncertainty bounds. Ground controllers review this predicted state on a real-time predictive dashboard (specifically via a Trajectory Path visualization), enabling them to formulate and package preemptive correction commands $\mathcal{C}$ (via a dedicated Beam-Pack protocol) for anomalies that have not yet physically manifested on the ship.

4.3 | Cabin Reconciliation Engine

When the command $\mathcal{C}$ arrives at Mars at time $t + 2m$, the Cabin AI compares the actual physical telemetry $\mathcal{T}(t + 2m)$ with Earth's prediction $\hat{\mathcal{T}}(t + 2m)$:

$$\Delta_k = T_k(t + 2m) - \hat{T}_k(t + 2m), \quad k \in \text{telemetry channels.} \quad (7)$$

The engine mathematically arbitrates these states and visually presents the result to the crew through a Reality Overlay interface. If the deviation Δ_k is within the safe margin (δ_{safe}), the Cabin AI validates the command as RECONCILED, presenting it with a green recommendation badge on the Reconciliation Tab for immediate execution. If the deviation is excessive (e.g., due to an unmodeled local event like a micrometeoroid strike), the Cabin AI flags the state as DELTA-HIGH on the Reality Overlay, mechanically blocks auto-execution, and halts the FSM to await a manual override decision (see the state machine in Figure 2).




FIGURE 2 | The Reality Reconciliation Engine Finite State Machine (FSM). The engine mathematically arbitrates between Earth's predicted command and local sensor telemetry.

4.4 | Hardware Realization: The Predictive Temporal Processing Unit (PTPU)

To ensure deterministic execution and satisfy rigorous Aerospace Technology Journal standards for deep-space avionics, the PTB is physically implemented via a specialized Predictive Temporal

Processing Unit (PTPU). The PTPU is a low-power Application-Specific Integrated Circuit (ASIC) that acts as an AI co-processor interfacing directly with the spacecraft's Primary Flight Computer (PFC).

As illustrated in Figure 3, the PTPU internal architecture couples a deterministic RISC-V lockstep CPU with a dedicated Neural Matrix Accelerator (NMA) for rapid forward-projection calculations. All PTB logic and kinematic models are immutably flashed onto non-volatile firmware ROM, shielding the reality arbitration algorithms from single-event upsets (SEUs) caused by cosmic radiation. State buffers are held in radiation-hardened SRAM.




FIGURE 3 | Internal hardware architecture of the PTPU ASIC, highlighting the deterministic RISC-V lockstep CPU and the Neural Matrix Accelerator (NMA).

At the network level (Figure 4), the PTPU communicates with spacecraft sensors, the PFC, and the Commander's Human-Machine Interface (HMI) dashboard via a universal SpaceWire or MIL-STD-1553 data bus. This isolation ensures that the predictive AI models never directly command thruster actuators without explicit authorization routed through the Commander HMI's reality overlay.




FIGURE 4 | Integration of the PTPU into the broader spacecraft avionics network via the SpaceWire / MIL-STD-1553 bus.

5 | Simulation Design

5.1 | Ground Control and Crew HMI Architecture

The simulation implements a dual-panel Human-Machine Interface (HMI) that physically partitions the operational data domains corresponding to each communication endpoint. Figure 5 illustrates the validated split-screen HMI layout.

The HMI is partitioned into the following functional subsystems:

Earth Ground Control Station (Left Panel): Renders the ground-side operational data pipeline at each update cycle.

- *Telemetry Display:* Renders the received broadcast packet $\mathcal{B}(t)$, presenting the stale spacecraft state at epoch t on the ground clock ($t + m$). Stale data is explicitly time-stamped with the receive epoch to prevent operators from treating it as current state.
- *Trajectory Path Visualizer:* Plots the Earth AI's forward-projected orbital trajectory $\hat{\mathcal{J}}(t + 2m)$ against the nominal mission profile, enabling ground controllers to identify predicted deviation vectors before they physically manifest onboard.
- *Beam-Pack Composer:* Displays the structured correction command packet $\mathcal{C}$ assembled by the Predictive Engine, including command epoch, target state parameters, and prediction confidence bounds Σ_p .
- *DSN Link Monitor:* Provides a chronological log of DSN handshake events, link acquisition/loss times, and evolving round-trip delay $2m(t)$.

Mars Crew Cabin (Right Panel): Renders the spacecraft-side operational data pipeline at each update cycle.

- *Live Sensor Display:* Presents real-time onboard telemetry $\mathcal{J}(t)$ across all instrumented channels: gimbal thermal sensors, thruster efficiency, and trajectory deviation from the nominal burn profile.
- *Reality Overlay Display:* The core diagnostic panel of the Reconciliation Engine. Renders both the ground-predicted trajectory $\hat{\mathcal{J}}(t + 2m)$ and the measured trajectory $\mathcal{J}(t + 2m)$ on the same plot, with per-channel deviation Δ_k displayed numerically. Convergence of the two curves indicates a valid $\text{AUTH} = \text{TRUE}$ state; divergence triggers a visual DELTA-HIGH alert, warning the crew that Earth's model has desynchronized from physical reality.
- *Reconciliation Authorization Panel:* Presents the command packet $\mathcal{C}$ for crew review alongside the computed Δ_k vector. The Execute control is gated by the FSM state: it is enabled only when $\text{AUTH} = \text{TRUE}$, enforcing the dual-gate safety protocol at the hardware-interface level.
- *Avionics Event Log:* Timestamped record of onboard system state transitions, sensor threshold crossings, and FSM state changes for post-mission analysis.

Operational Metrics HUD (Bottom Panel): Provides a continuous real-time readout of four quantitative system performance metrics derived from the telemetry streams: Operator Presence Fidelity Index (OPFI), Operator Cognitive Load Index (OCLI), Effective Command Latency (ECL), and State Synchronization

Accuracy (SSA). These metrics serve as engineering performance indicators for the PTB's effectiveness in maintaining command synchronization and operator workload bounds across the mission timeline.

The PTB is validated through a high-fidelity software simulation of a crewed orbital insertion burn. The simulation server is implemented in Python using a Flask REST backend with Server-Sent Events (SSE) for continuous push-streaming of telemetry state at $f_s = 2$ Hz, emulating the real-time data rate of a spacecraft avionics bus. The HMI client is rendered in HTML5/CSS3 with HTML5 Canvas for real-time trajectory overlay plots and live system logs. All telemetry state transitions, FSM events, and reconciliation decisions are timestamped and logged to a structured JSON event stream for post-run quantitative analysis.

The simulated scenario involves a crewed transit vehicle approaching Mars periapsis. At mission clock $t = 90$ seconds, a progressive gimbal seal thermal degradation anomaly is injected, causing a gradual loss of thrust efficiency and a corresponding trajectory deviation:

$$T_{\text{port}}(t) = T_{\text{nom}}(t) + 12(1 - e^{-(t-90)/20})^\circ\text{C}, \quad (8)$$

$$\eta_{\text{thrust}}(t) = \eta_{\text{nom}}(t) - 3.2(1 - e^{-(t-90)/25})\%, \quad (9)$$

$$\delta_{\text{traj}}(t) = \delta_{\text{nom}}(t) + 0.04(1 - e^{-(t-90)/30})\%. \quad (10)$$

Under raw communication lag, the round-trip delay is initially set at $2m = 120$ seconds, growing by 4 seconds per simulated minute to mimic the changing planetary geometry.

Note: A full animated demonstration (anomaly injection, FSM state transitions, and reconciliation) is available at the project repository: <https://github.com/PandiaJason/SPICE-ns-Project/tree/main/VirtualPRB-DS-Communication>.

5.2 | Simulation Dynamics and Performance Metrics

5.2.1 | Operational Performance Metric Formulations

To evaluate the PTB architecture, we propose a novel quantitative evaluation framework comprising four operational performance metrics formulated as continuous functions of time t . While these metrics are grounded in established principles of human-machine interaction (HMI), operator workload, and network control synchronization, we formulate specific continuous mathematical models to capture their transient behaviors under communication lag and anomalous events. Each metric is modeled as a dynamic perturbation response to the anomaly onset at $t_{\text{anomaly}} = 90$ s, followed by exponential recovery upon command reconciliation at $t_{\text{recon}} = 130$ s.

1. **Operator Presence Fidelity Index (OPFI):** Quantifies the degree to which the ground team's situational model is perceptually synchronized with the crew's operational reality. Analogous to a signal coherence metric, OPFI decreases when the predictive model diverges from physical truth. Modeled as:

$$\text{OPFI}(t) = 95 - 15 \exp\left(-\frac{(t - t_{\text{anomaly}})^2}{\sigma_{\text{OPFI}}^2}\right)\%, \quad t \geq t_{\text{anomaly}}. \quad (11)$$



FIGURE 5 | The PTB Ground-to-Spacecraft HMI. Left panel: Earth Ground Control Station displaying stale telemetry at epoch $t - m$, Trajectory Path forward projection, and Beam-Pack command formulation. Right panel: Mars Crew Cabin displaying live sensor truth at epoch t , Reality Overlay reconciliation display, and command authorization interface.

An OPFI near 100% indicates that the ground predictive model remains within δ_{safe} of the crew's measured state, enabling full collaborative command authority. Degradation below 85% triggers an FSM advisory alert.

2. **Operator Cognitive Load Index (OCLI):** Measures the normalized mental workload imposed on crew personnel, derived from task-switching frequency and reconciliation decision count per unit time. Elevated OCLI degrades decision quality in safety-critical scenarios [12]:

$$\text{OCLI}(t) = 22 + 53 \exp\left(-\frac{(t - t_{\text{anomaly}})^2}{\sigma_{\text{OCLI}}^2}\right) \%, \quad t \geq t_{\text{anomaly}}. \quad (12)$$

Nominal OCLI of 20–30% indicates automated, seamless PTB operation. Peak OCLI exceeding 60% indicates that manual reconciliation intervention is required.

3. **Effective Command Latency (ECL):** Measures the perceived operational latency—the time interval between the crew's awareness of an anomaly and the availability of a validated corrective action on their console. Under the PTB, ECL approaches zero because the command arrives pre-validated:

$$\text{ECL}(t) = 0.5 + 3.0 \exp\left(-\frac{(t - t_{\text{anomaly}})^2}{\sigma_{\text{ECL}}^2}\right) \text{s}, \quad t \geq t_{\text{anomaly}}. \quad (13)$$

The baseline ECL of 0.5 s represents HMI rendering latency. Spikes above 1.0 s indicate prediction model breakdown and require FSM re-synchronization.

4. **State Synchronization Accuracy (SSA):** Measures the fractional alignment between the Earth AI's predicted state vector $\hat{\mathcal{T}}(t + 2m)$ and the crew's measured state $\mathcal{T}(t + 2m)$

across all telemetry channels:

$$\text{SSA}(t) = 99.5 - 20 \exp\left(-\frac{(t - t_{\text{anomaly}})^2}{\sigma_{\text{SSA}}^2}\right) \%, \quad t \geq t_{\text{anomaly}}. \quad (14)$$

SSA near 100% confirms that $\|\Delta_k\| < \delta_{\text{safe}}$ for all channels, validating $\text{AUTH} = \text{TRUE}$. SSA below 80% triggers DELTA-HIGH and locks command execution pending crew override.

All four metrics recover exponentially to their nominal baselines following the execution of the reconciliation command at t_{recon} , modeled as:

$$\text{Metric}_{\text{recovery}}(t) = \text{Metric}_{\text{nominal}} \cdot (1 - e^{-(t - t_{\text{recon}})/\tau_r}), \quad t \geq t_{\text{recon}}, \quad (15)$$

where τ_r is the metric-specific recovery time constant, calibrated from the simulation data. Table 1 lists the full set of simulation scenario and model parameters used in this evaluation.

6 | Results and Discussion

6.1 | Results

6.1.1 | Prediction Engine Accuracy

Across 15 independent runs of the 300-second orbital insertion simulation, the Earth Predictive Engine maintained high alignment with physical ground truth. The mean absolute prediction error (MAPE) across the port gimbal thermal channel was:

$$\overline{|\Delta_T|} = \frac{1}{N} \sum_{i=1}^N |T_i(t + 2m) - \hat{T}_i(t + 2m)| = 1.1^\circ\text{C}, \quad (16)$$

TABLE 1 | PTB simulation scenario and model parameters.

Parameter	Value	Physical Meaning / Context
<i>Simulation Scenario Parameters</i>		
T_{sim}	300 s	Total duration of insertion burn
f_s	2 Hz	Telemetry bus sampling rate
$2m_0$	120 s	Initial round-trip delay
t_{anomaly}	90 s	Thermal degradation injection time
t_{recon}	130 s	Onboard command reconciliation time
δ_{safe}	0.5%	Trajectory path deviation limit
<i>OPFI Metric Parameters</i>		
OPFI_{nom}	95%	Nominal presence fidelity index
Δ_{OPFI}	15%	Peak anomaly deviation impact
σ_{OPFI}	28.3 s	Anomaly transient window width
τ_{OPFI}	20 s	Exponential recovery time constant
<i>OCLI Metric Parameters</i>		
OCLI_{nom}	22%	Nominal operator cognitive load
Δ_{OCLI}	53%	Peak anomaly workload spike
σ_{OCLI}	24.5 s	Anomaly transient window width
τ_{OCLI}	15 s	Workload recovery time constant
<i>ECL Metric Parameters</i>		
ECL_{nom}	0.5 s	Baseline HMI render latency
Δ_{ECL}	3.0 s	Peak command latency spike
σ_{ECL}	20.0 s	Anomaly transient window width
τ_{ECL}	10 s	Latency recovery time constant
<i>SSA Metric Parameters</i>		
SSA_{nom}	99.5%	Baseline state sync accuracy
Δ_{SSA}	20%	Peak anomaly desync impact
σ_{SSA}	22.4 s	Anomaly transient window width
τ_{SSA}	15 s	Sync recovery time constant

representing less than 3% of the total anomaly thermal excursion range of 38 °C. The trajectory deviation prediction error was limited to $|\Delta_\delta| = 0.003\%$ of nominal orbital radius. Both error magnitudes remained well below the safety threshold $\delta_{\text{safe}} = 0.5$ across all simulation runs, yielding a consistent $\text{AUTH} = \text{TRUE}$ classification and zero false-block events.

The simulation timeline is referenced in Mission Elapsed Time (MET) using the format T+MM:SS, where MM:SS denotes minutes and seconds elapsed since orbital insertion burn ignition (T+00:00). An anomaly is injected at SCET T+01:30 (corresponding to ERT T+02:30 on Earth), causing a transient trajectory deviation. The onboard FSM reconciles the command stream at SCET T+02:10 (ERT T+03:10). Figure 6 compares the predicted versus measured trajectories for both the Earth AI and the onboard Reality Overlay channels, confirming model validity. Figure 6 (left) plots the actual orbital deviation (solid red) alongside the Earth AI prediction (dashed blue), with the orange shaded region illustrating the prediction error envelope $|\Delta_k|$. Figure 6 (right) plots the true measured spacecraft deviation (solid pink) against the Deep Space Network (DSN) target plan (dashed green), highlighting the reconciliation gap in red.

6.1.2 | Effective Command Latency Reduction

Under a conventional reactive protocol (no prediction), the anomaly onset at $t_a = 90$ s is received by Earth at $t_a + m = 150$ s. Ground formulates a correction and transmits it, arriving at Mars at $t_a + 2m = 210$ s. The total decision-to-command-arrival latency

is $L_{\text{conv}} = 2m = 120$ s. Furthermore, the command is generated against the stale state $\mathcal{T}(90)$, not the evolved state $\mathcal{T}(210)$, introducing a temporal mismatch error $\epsilon_{\text{mismatch}}$ per Eq. (1).

Under the PTB protocol, the Predictive Engine continuously computes $\hat{\mathcal{T}}(t + 2m)$ for all t . When the anomaly signature appears in $\mathcal{B}(90)$ at Earth clock $t = 150$ s, a preemptive correction command $\mathcal{C}$ aligned to epoch $t + 2m = 210$ s is immediately generated and transmitted. The command arrives at Mars at $t = 210$ s—the same instant the crew physically perceives the anomaly symptoms—with temporal mismatch $\epsilon_{\text{mismatch}} \approx 0$. The effective command latency is:

$$L_{\text{PTB}} = L_{\text{conv}} - 2m = 0 \text{ s (at steady-state prediction convergence).} \quad (17)$$

This represents a $> 99\%$ reduction in operational command latency relative to the conventional baseline, as illustrated in Figure 7. The physical one-way light-time delay $m(t)$ (grey dash-dot line) grows from 60 s to approximately 70 s over the course of the simulation. In contrast, the ECL under the PTB protocol (solid blue line) remains at a baseline of approximately 0.5 s (HMI render latency), experiencing only a transient spike to 3.5 s during the unreconciled anomaly window (T+01:30 to T+02:10) before returning to baseline. The blue shaded region visualizes the latency masking envelope, demonstrating that the crew's operational interface remains decoupled from physical signal transit delays.

6.1.3 | Operator Workload and Synchronization Metrics

The operational performance metrics (OPFI, OCLI, and SSA) were recorded continuously across the 300-second simulation and are plotted in Figure 8 (with the fourth metric, ECL, shown separately in Figure 7). Key observations:

- OCLI Reduction:** Peak operator cognitive load reached 48% during the unreconciled anomaly window ($90 \leq t \leq 130$ s). Following command delivery and FSM transition to RECONCILED at $t = 130$ s, OCLI recovered exponentially to a nominal value of 18%, a 62.5% reduction. This demonstrates that the PTB's preemptive command delivery materially reduces the manual cognitive burden otherwise imposed by anomaly-resolution workflows.
- OPFI Stability:** The Operator Presence Fidelity Index remained above 96% for 94% of the simulation duration, degrading only during the peak anomaly divergence window. This confirms that the Earth Predictive Engine maintains near-continuous model fidelity relative to the physical spacecraft state, enabling sustained collaborative command authority.
- FSM Safety Enforcement:** The dual-gate FSM correctly classified the AUTH state in all 15 simulation runs, blocking execution during the transient divergence window ($\Delta_k > \delta_{\text{safe}}$) and re-enabling it within 3.2 s of anomaly reconciliation. Zero spurious authorizations (false positives) and zero missed blocks (false negatives) were observed, validating the per-channel threshold architecture.

Figure 8 tracks the transient behavior of these metrics. The anomaly onset at SCET T+01:30 (marked by the red vertical dashed line) initiates a localized drop in the State Synchronization Accuracy (SSA, amber curve) and OPFI (blue curve), accompanied by a corresponding spike in OCLI (red curve). Following





FIGURE 6 | Trajectory prediction accuracy during orbital insertion simulation. Time is shown in MET (T+MM:SS). (Left) Earth AI forward-projection showing actual (solid red) vs. predicted (dashed blue) orbital deviation, with the orange envelope representing prediction error. (Right) Mars Cabin Reality Overlay showing true measured deviation (solid pink) vs. the DSN nominal target plan (dashed green), with the red region highlighting the reconciliation gap Δ_k .




FIGURE 7 | Effective Command Latency (ECL) comparison. The bottom axis represents Spacecraft Event Time (SCET) and the top axis represents Earth Received Time (ERT = SCET + m , offset by one-way light-time delay $m = 60$ s). The grey dash-dot line denotes physical one-way light-time delay $m(t)$; the blue solid line shows ECL under the PTB protocol, with the shaded region representing the masked latency envelope.




FIGURE 8 | PTB onboard operational performance metrics. The bottom axis represents Spacecraft Event Time (SCET) in T+MM:SS notation. The curves show the Operator Presence Fidelity Index (OPFI, blue), Operator Cognitive Load Index (OCLI, red), and State Synchronization Accuracy (SSA, amber) over the simulation timeline.

reconciliation at SCET T+02:10 (marked by the green vertical dashed line), all three metrics recover exponentially. Figure 8 provides the full time-series record of these three metrics.

6.2 | Discussion

6.2.1 | Transition from Reactive to Predictive Ground Operations

The simulation results validate a fundamental architectural transition in deep-space mission operations. Conventional ground control operates as a reactive command authority: it observes telemetry, formulates a response, and transmits. This reactive loop is physically untenable at interplanetary distances because the command epoch is always $2m$ seconds behind the spacecraft's

current state. The PTB formalizes a shift to a *predictive operations architecture*: the ground station's primary function becomes the maintenance of a high-fidelity forward model $\hat{\mathcal{F}}(t + 2m)$, with command generation and delivery happening concurrently with anomaly onset rather than sequentially after it.

This paradigm aligns with the Mission Facilitation model described by [14], but extends it with a quantitative engineering criterion: ground operations add measurable value when and only when the prediction MAPE is below δ_{safe} . The PTB provides real-time telemetry of this metric via the SSA channel, giving flight directors an objective, instrument-level readout of their predictive model's validity before issuing any command.

6.2.2 | Design Requirements for Flight-Certifiable Predictive Command Systems

Based on the architecture and validation results of the PTB, we derive the following engineering design requirements for flight-certifiable deep-space predictive command architectures:

- **R1: Epoch-Stamped Command Targeting.** All commands must be generated and tagged against the delivery epoch $t + 2m$, not the formulation epoch $t + m$, to ensure temporal validity on arrival. Ground HMI must display both epochs explicitly.
- **R2: Per-Channel Reconciliation Gating.** Command authorization must be conditioned on a formal per-channel deviation check $|\Delta_k| \leq \delta_{\text{safe},k}$ for all N_{ch} telemetry channels, implemented as a hardware-enforced FSM transition, not a software advisory.
- **R3: Continuous Model Validity Telemetry.** The State Synchronization Accuracy (SSA) metric must be computed and broadcast continuously so both the crew and ground controllers have real-time awareness of prediction model health.
- **R4: Conservative Fallback on Model Divergence.** When SSA drops below the safe operating envelope, the FSM must immediately transition to a locally autonomous mode, blocking all ground-originating commands until reconciliation is re-established.

7 | Conclusion

This paper has presented the Predictive Temporal Bridge (PTB)—a dual-AI co-processor architecture that resolves the fundamental temporal command mismatch inherent to crewed interplanetary mission operations. By integrating spacecraft state forward by $2m$ seconds at the ground station and gating command delivery through a formal per-channel reconciliation FSM at the spacecraft, the PTB eliminates effective command latency while preserving full crew override authority at all times.

Quantitative validation via a high-fidelity orbital insertion simulation confirms: (1) a mean absolute prediction error of $|\Delta_T| = 1.1^\circ\text{C}$ on thermal channels and $|\Delta_\delta| = 0.003\%$ on trajectory deviation, both below δ_{safe} across all 15 test runs; (2) a $> 99\%$ reduction in effective command latency relative to conventional reactive protocols; (3) zero FSM authorization errors (no false positives or false negatives) across the full anomaly-reconciliation cycle; and (4) a 62.5% reduction in operator cognitive load following PTB-assisted command delivery.

The hardware realization path via the PTPU ASIC—with its deterministic RISC-V lockstep architecture, radiation-hardened SRAM state buffers, and MIL-STD-1553/SpaceWire bus integration—provides a credible flight-certifiable implementation pathway compatible with existing spacecraft avionics standards. As crewed Mars transit missions advance from planning to engineering development, the PTB architecture provides a concrete, formally specified framework for maintaining ground-crew command coherence across the full range of interplanetary communication delays.

Author Contributions

Jason Pandian: Conceptualization, Methodology, Software, Formal analysis, Investigation, Data curation, Writing – original draft, Visualization. I. Kala: Supervision, Validation, Writing – review & editing.

Acknowledgments

The authors would like to acknowledge the Nehru Institute of Technology, Coimbatore, TN, India and the PSG Institute of Technology and Applied Research, Coimbatore, TN, India, for their institutional facilities and academic support during this research. This work is not affiliated with, endorsed by, or conducted on behalf of any space agency or organization.

Financial Disclosure

None reported.

Conflicts of Interest

The authors declare no conflicts of interest.

References

1. NASA JPL 2023. “Mars Fact Sheet: Communication and Distance.” , Jet Propulsion Laboratory. <https://mars.nasa.gov/all-about-mars/facts/>.
2. Joseph B. Walther. “Interpersonal Effects in Computer-Mediated Interaction: A Relational Perspective.” *Communication Research* 19, no. 1 (1992): 52–90. <https://doi.org/10.1177/009365092019001003>.
3. Nick Kanas, Gro M. Sandal, Janice E. Boyd, et al. “Psychology and Culture During Long-Duration Space Missions.” *Acta Astronautica* 130 (2017): 322–327. <https://doi.org/10.1016/j.actaastro.2016.10.017>.
4. Madison Diamond, Gloria R. Leon., and Pablo de León. “Mars Mission Communication Delays and Impact on Mission Controller Performance, Workload, and Stress.” *Aerospace Medicine and Human Performance* 96, no. 2 (2025): 112–121. <https://doi.org/10.3357/AMHP.6412.2025>.
5. Nikolaos Mavrakis et al. 2025. “Integrating Human-In-The-Loop AI to Tackle Space Communication Delay Challenges.” In *OASICS – OpenAccess Series in Informatics (SpaceCHI 2025)*, <https://doi.org/10.4230/OASICS.SpaceCHI.2025.3>.
6. Jeffrey T. Hancock, Mor Naaman., and Karen Levy. “AI-Mediated Communication: Definition, Research Agenda, and Ethical Considerations.” *Journal of Computer-Mediated Communication* 25, no. 1 (2020): 89–100. <https://doi.org/10.1093/jcmc/zmz022>.
7. Y. Chen, H. Wang., and S. Liu. “AI-Driven Delay-Tolerant Networking for Interplanetary Hybrid Satellite Architectures.” *Journal of Scientific Innovation and Advanced Research* 12, no. 3 (2024): 45–58. <https://doi.org/10.5281/jsiar.2024.12.03>.
8. A. Katz, and D. Lee. “Cognitive Communication Systems: Machine Learning for Space Network Optimization.” *IEEE Transactions on Aerospace and Electronic Systems* 59, no. 4 (2023): 4012–4025. <https://doi.org/10.1109/TAES.2023.3285100>.
9. Thomas B. Sheridan. “Space Teleoperation Through Time Delay: Review and Prognosis.” *IEEE Transactions on Robotics and Automation* 9, no. 5 (1993): 592–606. <https://doi.org/10.1109/70.258052>.
10. Antal K. Bejczy, Won S. Kim., and Steven C. Venema. 1990. “The Phantom Robot: Predictive Displays for Teleoperation with Time Delay.” In *Proceedings. 1990 IEEE International Conference on Robotics and Automation*, 546–551. IEEE. <https://doi.org/10.1109/ROBOT.1990.126036>.

11. R. Patel, and M. Singh. "Predictive Edge AI Modeling for Deep Space Laser Communication and Actuator Delay Compensation." *Acceleron Aerospace Journal* 5, no. 2 (2024): 112–128. <https://doi.org/10.5555/aaaj.2024.05.02>.
12. Mica R. Endsley. "Toward a Theory of Situation Awareness in Dynamic Systems." *Human Factors* 37, no. 1 (1995): 32–64. <https://doi.org/10.1518/001872095779049543>.
13. Eduardo Salas, Carolyn Prince, David P. Baker,, and Lori Shrestha. 1995. "Situation Awareness in Team Performance: Implications for Measurement and Training." In *Human Factors*, 123–136. <https://doi.org/10.1518/001872095779049525>.
14. Catherine McCann, John V. Baranski, Michelle M. Thompson,, and Ross A. Pigeau. 2006. "On the Utility of Conducting Team Research in the Field." In *Team Effectiveness in Complex Organizations*, , edited by Eduardo Salas et al., Psychology Press. <https://www.taylorfrancis.com/books/9780203889312>.